



Advanced Data Science

Lecture 5 : Introduction to Statistical Learning

Carl Henrik Ek - che29@cam.ac.uk

17th of November, 2021

<http://carlhenrik.com>



Access how to get and combine the data-sources for a potential problem

Access how to get and combine the data-sources for a potential problem

Assess things to do with data before you have a question

Access how to get and combine the data-sources for a potential problem

Assess things to do with data before you have a question

Address what you can do when you have data and a question to answer

Access

- ?

Assess

- Introduction to Probability (IA)
- Scientific Computing (IA)
- Cloud Computing (II)
-

Address

- ML and Real-world Data (IA)
- Data Science (IB)
- AI (IB)
- ML & Bayesian Inference (II)
- Deep NN (II)
- Randomised Algorithms (II)
-

Why?




*You need to put Machine Learning in the **context** of data (and humans)*
– Neil Lawrence

- Tasks that are too hard to program
 - speech recognition
 - image understanding

- Tasks that are too hard to program
 - speech recognition
 - image understanding
- Tasks beyond our capability
 - weather prediction
 - web search

- Tasks that are too hard to program
 - speech recognition
 - image understanding
- Tasks beyond our capability
 - weather prediction
 - web search
- Machine Learning bridges the knowledge gap by data

Task of Statistical Learning



$$p(\theta \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \theta)p(\theta)}{p(\mathcal{D})}$$

- Inductive biases comes into the learning procedure

$$p(\theta \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \theta)p(\theta)}{p(\mathcal{D})}$$

- Inductive biases comes into the learning procedure
- *Most knowledge is introduced **before** we apply ML*
 - access** what data did I acquire?
 - assess** how did I prepare/treat the data?

$$p(\theta \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \theta)p(\theta)}{p(\mathcal{D})}$$


- Inductive biases comes into the learning procedure
- *Most knowledge is introduced **before** we apply ML*
 - access** what data did I acquire?
 - assess** how did I prepare/treat the data?
- The idea of the 80/20

- what can machine learning actually do?
- what role does the "data scientist" play in the machine learning loop?
- put machine learning into context

Statistical Learning

Learning is the process of converting experience into expertise or knowledge.

– Sha Ben-David

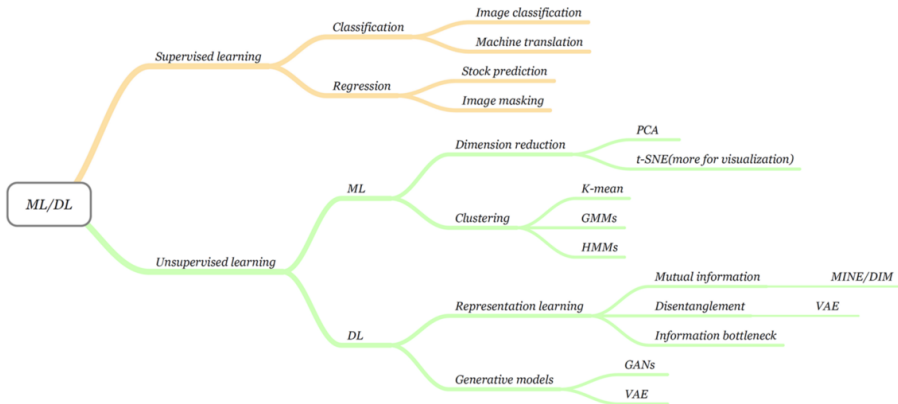


Supervised Learning $p(y \mid x)$

"Unsupervised" Learning $p(y)$

Reinforcement Learning $p(\pi, f \mid \mathcal{L})$

Machine Learning Methods



Domain Set \mathcal{X} the set of measurements/objects that we want to label (input)

Domain Set \mathcal{X} the set of measurements/objects that we want to label (input)

Label Set \mathcal{Y} the set of outputs

Domain Set \mathcal{X} the set of measurements/objects that we want to label (input)

Label Set \mathcal{Y} the set of outputs

Training Data \mathcal{S} a finite sequence of pairs in $\mathcal{X} \times \mathcal{Y}$

Data Distribution \mathcal{D} probability distribution governing the measurements


Data Distribution \mathcal{D} probability distribution governing the measurements

Data Generation $f : \mathcal{X} \rightarrow \mathcal{Y}$ the underlying generating process that we wish to recover

Data Distribution \mathcal{D} probability distribution governing the measurements

Data Generation $f : \mathcal{X} \rightarrow \mathcal{Y}$ the underlying generating process that we wish to recover

Prediction Rule $h : \mathcal{X} \rightarrow \mathcal{Y}$ what we wish to recover, the object that encodes the recovered knowledge



$$L_{\mathcal{D},f}(h) := \mathcal{D}(\{x : h(x) \neq f(x)\})$$

- measure of success as probability of misclassified points (true risk)

$$L_{\mathcal{D},f}(h) := \mathcal{D}(\{x : h(x) \neq f(x)\})$$


- measure of success as probability of misclassified points (true risk)
- we do not have access to \mathcal{D}

$$L_{\mathcal{D},f}(h) := \mathcal{D}(\{x : h(x) \neq f(x)\})$$

- measure of success as probability of misclassified points (true risk)
- we do not have access to \mathcal{D}
- we do not have access to f

$$L_{\mathcal{S}}(h) := \frac{|\{i \in [m] : h(x_i) \neq y_i\}|}{m}$$

- We **assume** that $\mathcal{S} \sim \mathcal{D}$
- Empirical measure of risk




$$\mathcal{D} = \frac{1}{3}\mathcal{N}(\cdot, \cdot) + \frac{2}{3}\mathcal{N}(\cdot, \cdot)$$

$$h_{\mathcal{S}}(x) = \begin{cases} y_i & \text{if } \exists i \in [m] \text{ s.t. } x_i = x \\ 0 & \text{otherwise} \end{cases}$$

- $L_{\mathcal{S}}(h_{\mathcal{S}}) = 0$ for all training data-sets
- if label 0 corresponds to **red**
 $L_{\mathcal{D}}(h_{\mathcal{S}}) = \frac{1}{3}$
- if label 0 corresponds to **blue**
 $L_{\mathcal{D}}(h_{\mathcal{S}}) = \frac{2}{3}$

True Risk is a Random Variable



$$L_{\mathcal{S}}(A(\mathcal{S})) := \frac{|\{i \in [m] : h(x_i) \neq y_i\}|}{m}$$

- We use an algorithm $A : \mathcal{S} \rightarrow h$ to find a hypothesis

$$h_S \in \operatorname{argmin}_{h \in \mathcal{H}} L_S(h)$$

- We cannot parametrise **all** possible hypothesis

m How much data do I need?

m How much data do I need?

A How much does my solution depend on what I find?

m How much data do I need?

A How much does my solution depend on what I find?

H How does my solution depend on the hypothesis class I choose?

- We will assume the probability to get a non-representative sample to be δ

- We will assume the probability to get a non-representative sample to be δ
 - $(1 - \delta)$ is the confidence in our prediction

- We will assume the probability to get a non-representative sample to be δ
 - $(1 - \delta)$ is the confidence in our prediction
- We have a threshold ϵ where we accept the hypothesis

$$L_{\mathcal{D},f}(h_{\mathcal{S}}) \leq \epsilon$$

- We will assume the probability to get a non-representative sample to be δ
 - $(1 - \delta)$ is the confidence in our prediction
- We have a threshold ϵ where we accept the hypothesis

$$L_{\mathcal{D},f}(h_{\mathcal{S}}) \leq \epsilon$$

- We are looking for a *hypothesis* that will be **probably** (with confidence $(1 - \delta)$) **approximately** (up to an error ϵ) **correct**.


$$\mathcal{S} \sim \mathcal{D}^m$$

- each sample is *independent* from the other
- the *order* of samples does not effect the data distribution
- the sampling process does not effect the "world"

$$m \geq \frac{\log \left(\frac{|\mathcal{H}|}{\delta} \right)}{\epsilon}$$

- PAC learning allows us to provide bounds on the learning procedure
- How much data do we need?
- How large hypothesis class can we allow?

The Bias-Complexity Trade-off




h^* the optimal predictor

h_{opt} the optimal hypothesis

\hat{h}_{opt} the optimal hypothesis on training data


\hat{h} the hypothesis found by learning algorithm

The Bias-Complexity Trade-off



$$\begin{aligned} & \epsilon(\hat{h}) - \epsilon(h^*) \\ &= \underbrace{\epsilon(h_{\text{opt}}) - \epsilon(h^*)}_{\text{Approximation}} \\ &+ \underbrace{\epsilon(\hat{h}_{\text{opt}}) - \epsilon(h_{\text{opt}})}_{\text{Estimation}} \\ &+ \underbrace{\epsilon(\hat{h}) - \epsilon(\hat{h}_{\text{opt}})}_{\text{Optimisation}} \end{aligned}$$

The Bias-Complexity Trade-off



High Complexity low bias (ϵ_{app} small), but high risk of overfitting (ϵ_{est} large)

Low Complexity high bias (ϵ_{app} large), low risk of overfitting (ϵ_{est} small)



Theorem (The No-Free-Lunch Theorem)


Let A be any learning algorithm for the task of binary classification with respect to 0 – 1 loss over the domain \mathcal{X} . Let m be any number smaller than $\frac{|\mathcal{X}|}{2}$. Then there exists a distribution $\mathcal{D}(\{\mathcal{X} \times \{0, 1\}\})$ such that,

- *There exists a function $f : \mathcal{X} \rightarrow \{0, 1\}$ with $L_{\mathcal{D}}(f) = 0$*
- *With probability at least $\frac{1}{7}$ over the choice of $S \sim \mathcal{D}^m$ we have $L_{\mathcal{D}}(A(S)) \geq \frac{1}{8}$*


- There exists no universal learner
- For every learner there exist a task on which it fails
- Every algorithm that learns something useful does so by assumptions

- There exists no universal learner
- For every learner there exist a task on which it fails
- Every algorithm that learns something useful does so by assumptions
- There is no free lunch algorithm

- We can never have sufficient data
- We can never find a method that will guarantee to find the right solution
- We can never be certain about the true risk of our outcome



- Shai Shalev-Shwartz et al. (2014). *Understanding Machine Learning: From Theory to Algorithms*. New York, NY, USA: Cambridge University Press
<https://www.cs.huji.ac.il/~shais/UnderstandingMachineLearning/>




- O. Bousquet et al. (2004). “Introduction to Statistical Learning Theory”. In: vol. Lecture Notes in Artificial Intelligence 3176. Heidelberg, Germany: Springer, pp. 169–207,
http://www.econ.upf.edu/~lugosi/mlss_sl_t.pdf



Explicit vs. Tacit Knowledge



The Gap



Access enormous inductive bias in what data to acquire

Access enormous inductive bias in what data to acquire


Assess human bias in what questions will probably be asked


Access enormous inductive bias in what data to acquire

Assess human bias in what questions will probably be asked

Address "it is just curve fitting"

Curve Fitting is Really Fun





Today/Mon Generalised Linear Models

Today/Mon Generalised Linear Models

Wen Unsupervised Learning


Today/Mon Generalised Linear Models

Wen Unsupervised Learning

Fri Approximate Inference

Generalised Linear Models


$$h \in \mathcal{H}$$



$\mathbf{x} \in \mathcal{X}$ explanatory variable

$y \in \mathcal{Y}$ response variable

Task *explain the response by the explanatory variables*



$$y_i = \sum_{j=1}^d \beta_j x_{ij} + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2).$$

$$\mathbb{E}[y_i \mid \mathbf{x}_i] = \mathbb{E} \left[\sum_{j=1}^d \beta_j x_{ij} + \epsilon \right]$$

$$\begin{aligned}\mathbb{E}[y_i \mid \mathbf{x}_i] &= \mathbb{E} \left[\sum_{j=1}^d \beta_j x_{ij} + \epsilon \right] \\ &= \mathbb{E} \left[\sum_{j=1}^d \beta_j x_{ij} \right] + \mathbb{E}[\epsilon]\end{aligned}$$

$$\begin{aligned}\mathbb{E}[y_i \mid \mathbf{x}_i] &= \mathbb{E} \left[\sum_{j=1}^d \beta_j x_{ij} + \epsilon \right] \\ &= \mathbb{E} \left[\sum_{j=1}^d \beta_j x_{ij} \right] + \mathbb{E}[\epsilon] \\ &= \sum_{j=1}^d \beta_j x_{ij} + 0.\end{aligned}$$

$$y_i = \sum_{j=1}^d \beta_j x_{ij} + \epsilon,$$

$$y_i = \sum_{j=1}^d \beta_j x_{ij} + \epsilon,$$

$$y_i + \epsilon = \sum_{j=1}^d \beta_j x_{ij},$$

$$y_i = \sum_{j=1}^d \beta_j x_{ij} + \epsilon,$$

$$y_i + \epsilon = \sum_{j=1}^d \beta_j x_{ij},$$

$$\hat{y}_i = \sum_{j=1}^d \beta_j x_{ij},$$

$$y_i = \sum_{j=1}^d \beta_j x_{ij} + \epsilon,$$

$$y_i + \epsilon = \sum_{j=1}^d \beta_j x_{ij},$$

$$\hat{y}_i = \sum_{j=1}^d \beta_j x_{ij},$$

$$\hat{y}_i \sim \mathcal{N}(y_i, \sigma^2) = \mathcal{N} \left(\sum_{j=1}^d \beta_j x_{ij}, \sigma^2 \right),$$

$$g(\mathbb{E}[y_i \mid \mathbf{x}_i]) = \sum_{j=1}^d \beta_j x_{ij},$$

$g(\cdot)$ link function

$y \sim \mathcal{D}$ Exponential Dispersion Family

$\sum_{j=1}^d \beta_j x_{ij}$ Linear predictor

$$\mathbb{E}[y_i \mid \mathbf{x}_i] = g^{-1}\left(\sum_{j=1}^d \beta_j x_{ij}\right),$$

- The inverse of the *link* maps the linear predictor to the first moment of the response
- Linear regression the link is identity

¹[https://towardsdatascience.com/
glms-part-iii-deep-neural-networks-as-recursive-generalized-linear-URL](https://towardsdatascience.com/glms-part-iii-deep-neural-networks-as-recursive-generalized-linear-URL)

$$\mathbb{E}[y_i \mid \mathbf{x}_i] = g^{-1}\left(\sum_{j=1}^d \beta_j x_{ij}\right),$$

- The inverse of the *link* maps the linear predictor to the first moment of the response
- Linear regression the link is identity
- Looks an awful lot like a neural network¹

¹<https://towardsdatascience.com/>

$$f(y; \theta, \phi) = e^{\frac{\theta y - b(\theta)}{a(\phi)} + c(y, \phi)},$$

θ location parameter

ϕ scale parameter

²https://en.wikipedia.org/wiki/Exponential_dispersion_model

$$f(y; \mu, \sigma^2) = e^{\frac{\mu y - \frac{1}{2}\mu^2}{\sigma^2} - \frac{y^2}{2\sigma^2} - \frac{1}{2} \ln(2\pi\sigma^2)}$$

- $\theta = \mu$
- $\phi = \sigma^2$
- $b(\theta) = \frac{1}{2}\mu^2$
- $a(\phi) = \sigma^2$
- $c(y, \phi) = \frac{y^2}{2\sigma^2} - \frac{1}{2} \ln(2\pi\sigma^2)$

$$\mathbb{E}[y \mid \mathbf{x}] = \frac{\partial}{\partial \theta} b(\theta)$$
$$\mathbb{V}[y \mid \mathbf{x}] = a(\phi) \frac{\partial^2}{\partial \theta^2} b(\theta).$$

- Through a consistent parametrisation we can generalise the moment calculations

Model	Response Variable	Link	Explanatory Variable
Linear Regression	Normal	Identity	Continuous
Logistic Regression	Binomial	Logit	Mixed
Poisson Regression	Poisson	Log	Mixed
ANOVA	Normal	Identity	Categorical
ANCOVA	Normal	Identity	Mixed
Loglinear	Poisson	Log	Categorical
Multinomial response	Multinomial	Generalized Logit	Mixed

Summary

- Brief introduction to statistical learning theory
- Take home
 - ML models and algorithms is only a small part of the story
 - we are doing a lot better than we should be
 - we are not sure what we are doing but it somehow works

- Brief introduction to statistical learning theory
- Take home
 - ML models and algorithms is only a small part of the story
 - we are doing a lot better than we should be
 - we are not sure what we are doing but it somehow works
 - it is not explicit knowledge that pushes data-science forward, it is tacit and implicit

- Generalised Linear Models
 - "main" tool for statisticians
 - very well studied, excellent literature
 - linearity provides means of interpretability
 - excellent software packages `statsmodels`

Monday (22/11) Lecture: Generalised Linear Models





Tuesday (23/11) Lab: Generalised Linear Models

Wednesday (24/11) Lecture: Unsupervised Learning/Visualisation

Thursday (25/11) Tick: Generalised Linear Models

Friday (26/11) Lecture: Statistical Inference

eof

-  Bishop, Christopher M. (2006). *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Secaucus, NJ, USA: Springer-Verlag New York, Inc.
-  Bousquet, O., S. Boucheron, and G. Lugosi (2004). “Introduction to Statistical Learning Theory”. In: vol. Lecture Notes in Artificial Intelligence 3176. Heidelberg, Germany: Springer, pp. 169–207.
-  McCullagh, P. and J. A. Nelder (1989). *Generalized Linear Models*. London, UK: Chapman Hall / CRC: Chapman Hall / CRC.
-  Shalev-Shwartz, Shai and Shai Ben-David (2014). *Understanding Machine Learning: From Theory to Algorithms*. New York, NY, USA: Cambridge University Press.